

String Matching


Reads are strings


GTATGCACGCGATAG	TATGTCGCAGTATCT	CACCCTATGTCGCAG
TAGCATTGCGAGACG	GGTATGCACGCGATA	TGGAGCCGGAGCACC
TGTCTTTGATTCCTG	CGCGATAGCATTGCG	GCATTGCGAGACGCT
GACGCTGGAGCCGGA	GCACCCTATGTCGCA	GTATCTGTCTTTGAT
TATCGCACCTACGTT	CAATATTCGATCATG	GATCACAGGTCTATC
CACGGGAGCTCTCCA	TGCATTTGGTATTTT	CGTCTGGGGGGGTATG
GTATGCACGCGATAG	ACCTACGTTCAATAT	TATTTATCGCACCTA
GCGAGACGCTGGAGC	CTATCACCCCTATTAA	CTGTCTTTGATTCCT
CCTACGTTCAATATT	GCACCTACGTTCAAT	GTCTGGGGGGGTATGC
GACGCTGGAGCCGGA	GCACCCTATGTCGCA	GTATCTGTCTTTGAT
TATCGCACCTACGTT	CAATATTCGATCATG	GATCACAGGTCTATC
CACGGGAGCTCTCCA	TGCATTTGGTATTTT	CGTCTGGGGGGGTATG

Sequencing reads are *strings*; sequences of characters


The strings are the only hints we get about *where the reads came from* with respect to the longer DNA molecules...

... like pictures on puzzle pieces





What if I told you to find all the places where the string GATACCA occurs in here?



What if I told you to find all the places where the string
GATACCA occurs in here?

Strings

Read

CTCAAACCTCTGACCTTTGGTGATCCACCCGCCTAGGCCTTC x billions

Reference

GATCACAGGTCTATCACCTATTAACCACTCACGGGAGCTCTCCATGCATTTGGTATTTT
CGTCTGGGGGGTATGCACGCGATAGCATTGCGAGACGCTGGAGCCGGAGCACCTATGTC
GCAGTATCTGTCTTTGATTCTGCCTCATCCTATTATTTATCGCACCTACGTTCAATATT
ACAGGCGAACATACTTACTAAAGTGTGTTAATTAATTAATGCTTGTAGGACATAATAATA
ACAATTGAATGTCTGCACAGCCACTTTCCACACAGACATATAACAAAAAATTTCCACCA
AACCCCCCTCCCCGCTTCTGGCCACAGCACTTAAACAGCTCTGCCAAACCCCAAAA
ACAAAGAACCCTAACACCAGCCTAACCACTTTCAAATTTTATCTTTGGCGGTATGCAC
TTTTAACAGTCACCCCCCACTAACATTTATTTTCCCCTCCCACTTCCATACTACTAAT
CTCATCAATACAACCCCCGCCATCTTACCAGCACACACACACCCCTTAAACCCATA
CCCCGAACCAACCAAAACCCCAAAAGTACACCCCCACAGTTTATGTAGCTTACCTCCTCAA
GCAATACACTGACCCGCTCAAACCTCTGGATTTTGGATCCACCCAGCGCTTTGGCCTAAA
CTAGCCTTTCTATTAGCTCTTAGAAGATTACACATGCAAGCATCCCCGTCCAGTGAGT
TCACCCTCTAAATCACCACGATCAAGGAACAAGCATCAAGCACGCAGTAATGCAGCTC
AAAACGCTTAGCCTAGCCACACCCTCACGGGAAACAGCAGTGATTAACCTTAGCAATAA
ACGAAAGTTTAACTAAGCTATACTTACCCAGGGTTGGTCAATTTCTGTCAGCCACCGC
GGTCACACGATTAACCCAAGTCAATGAAGCCGGCGTAAAGAGTGTTTAGATCACCCCC
TCCCCAATAAAGCTAAAACTCACCTGATTTGTAAAAAATCCAGTTTACAAAAATAGAC
TACGAAAGTGGCTTTAACATATCTGAACAACAATAGCTAAGATCTGGGATTAGA
TACCCCACTATGCTTAGCCCTAAACCTCAACAGTAAACAAATAGGCGCCAGAA
CACTACGAGCCACAGCTTAAAACTCAAAGGACCTGGCGGTGCTTCATTTAGAGG
AGCCTGTTCTGTAATCGATAAACCCGATCAACCTCACACCTCTTGCTTTATATA
CCGCCATCTTCAGCAAACCTGATGAAGGCTACAAAGTAAGCGCAAGTACCTAG
ACGTTAGGTCAAGGTGTAGCCCATGAGGTGGCAAGAAATGGGCTACATTTTC
AAAACCTACGATAGCCCTTATGAACTTAAGGGTCAAGGTGGATTTAGCAGTAA
AGTAGAGTGCTTAGTTGAACAGGGCCCTGAAGCGCGTACACACCCGCCCGTCACCC
AAGTATACTTCAAAGGACATTTAACTAAAACCCCTACGCATTTATATAGAGGAGACA
CGTAACCTCAAACCTCTGCCTTTGGTGATCCACCCGCCTTGGCCTACCTGCATAATGAAG
AAGCACCCAACCTTAACTTAGGAGATTTCACTTAACTTGACCGCTCTGAGCTAAACCTA
GCCCCAAACCCACTCCACCTTACTACCAGACAACCTTAGCCAAACCATTTACCCAAATAA
AGTATAGGCGATAGAAATTGAAACCTGGCGCAATAGATATAGTACCGCAAGGGAAAGATG
AAAAATTATAACCAAGCATAATATAGCAAGGACTAACCCCTATACCTTCTGCATAATGAA
TTAACTAGAAATAACTTTGCAAGGAGAGCCAAAGCTAAGACCCCCGAAACCAGACGAGCT
ACCTAAGAACAGCTAAAAGAGCACACCCGTCTATGTAGCAAAATAGTGGGAAGATTTATA
GGTAGAGGCGACAAACCTACCGAGCCTGGTGATAGCTGGTTGTCCAAGATAGAATCTTAG
TTCAACTTTAAATTTGCCCACAGAACCCTCTAAATCCCCTTGTAATTTAACTGTTAGTC
CAAAGAGGAACAGCTCTTTGGACACTAGGAAAAAACCTTGTAAGAGAGGTAAAAAATTTA

x million

We're going to *need* the
right algorithms...

Strings are well studied

Many kinds of data are string-like: books, web pages, files on your hard drive, medical records, chess games, ...

Algorithms for one kind of string are often applicable to others:

Regular expression matching can find files on your filesystem (grep), or bad network packets (snort)

Indexes for books and web pages (inverted indexing) can be used to index DNA sequences

Methods for understanding speech (HMMs) can be used to understand handwriting or identify genes in genomes

Strings come from somewhere

Processes that give rise to real-world strings are complicated. It helps to understand them.

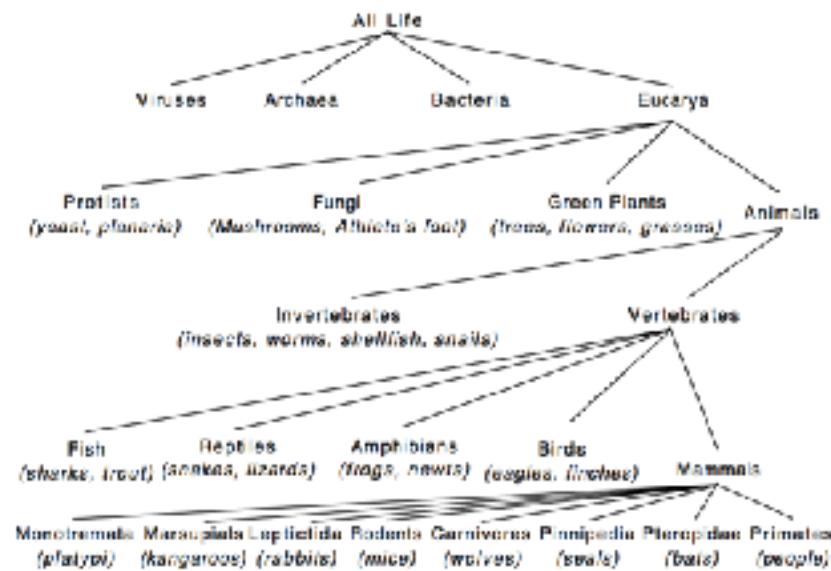



Figure from: Hunter, Lawrence. "Molecular biology for computer scientists." *Artificial intelligence and molecular biology* (1993): 1-46.

1. Evolution: Mutation
Recombination
(Retro)transposition



2. Lab procedures: PCR
Cell line passages



3. Sequencing: Fragmentation bias
Miscalled bases

Strings have structure

One way to model a string-generating process is with coin flips:

$$\{ \text{heads, heads} = A, \text{heads, tails} = C, \text{tails, heads} = G, \text{tails, tails} = T \}$$

But such strings lack internal patterns (“structure”) exhibited by real strings

> 40% of the human genome is covered by *transposable elements*, which copy-and-paste themselves across the genome and mutate




Image from: Cordaux R, Batzer MA. The impact of retrotransposons on human genome evolution. Nat Rev Genet. 2009 Oct;10(10):691-703

Slipped strand mis-pairing during DNA replication results in expansion or retraction of simple (*tandem*) repeats

... **ATATATATATAT** ...
↕
... **ATATATATATATATATAT** ...

String definitions

String S is a finite sequence of characters

Characters are drawn from alphabet Σ

Usually, $\Sigma = \{ A, C, G, T \}$

$|S|$ = number of characters in S

```
>>> s = 'ACGT'  
>>> len(s)  
4
```

ε is “empty string” $|\varepsilon| = 0$

```
>>> len('')  
0
```

String definitions

Positions within a string *S* are referred to with *offsets*

```
>>> s = 'ACGT'  
>>> s[0]  
'A'  
>>> s[2]  
'G'
```

Leftmost offset = 0 in Python and most other languages

String definitions

Concatenation of S and T , ST = characters of S followed by characters of T

```
>>> s = 'AACCC'  
>>> t = 'GGTT'  
>>> s + t  
'AACCCGGTT'
```


String definitions

Substring of S is a string occurring inside S

```
>>> s = 'AACCGGTT'  
>>> s[2:6]  
'CCGG' # substring of seq
```

S is a *substring* of T if there exist (possibly empty) strings u and v such that $T = uSv$

String definitions

Prefix of S is a substring starting at the beginning of S

```
>>> s = 'AACCGGTT'  
>>> s[0:6]  
'AACCGG' # prefix  
>>> s[:6] # same as above  
'AACCGG'
```

S is a *prefix* of T if there exists a string u such that...

String definitions

Prefix of S is a substring starting at the beginning of S

```
>>> s = 'AACCGGTT'  
>>> s[0:6]  
'AACCGG' # prefix  
>>> s[:6] # same as above  
'AACCGG'
```

S is a *prefix* of T if there exists a string u such that $T = Su$

String definitions

Suffix is substring ending at end of S

```
>>> s = 'AACCGGTT'
>>> s[4:8]
'GGTT' # suffix
>>> s[4:] # like s[4:len(s)]
'GGTT'
>>> s[-4:] # like s[len(s)-4:len(s)]
'GGTT'
```

S is a *suffix* of T if there exists a string u such that...

String definitions

Suffix is substring ending at end of S

```
>>> s = 'AACCGGTT'
>>> s[4:8]
'GGTT' # suffix
>>> s[4:] # like s[4:len(s)]
'GGTT'
>>> s[-4:] # like s[len(s)-4:len(s)]
'GGTT'
```

S is a *suffix* of T if there exists a string u such that $T = uS$

String definitions

Usually assume alphabet Σ is finite, with $O(1)$ elements

Nucleic acid alphabet: $\{A, C, G, T\}$

Amino acid alphabet: $\{A, R, N, D, C, E, Q, G, H, I, L, K, M, F, P, S, T, W, Y, V\}$

Occasionally we'll consider what happens as $|\Sigma|$ grows

Exact matching

Find places where *pattern* P occurs as a substring of *text* T .
Each such place is an *occurrence* or *match*.

Let $n = |P|$, and let $m = |T|$ Assume $n \leq m$

Alignment: a way of putting P 's characters opposite T 's.
May or may not correspond to an match.

P : word

T : There would have been a time for such a word

Alignment 1: word

Alignment 2: word

Exact matching

What's a simple algorithm for exact matching?

P: word

T: There would have been a time for such a word

word word word word word word word word word word **word**

word word word word word word word word word

word word word word word word word word word

word word word word word word word word word

word word word word word word word word word

One
occurrence



Try all possible alignments. For each, check if it matches.
This is the *naïve algorithm*.

Exact matching: naïve algorithm

```
def naive(p, t):  
    occurrences = []  
    for i in range(len(t) - len(p) + 1):  # Loop over alignments  
        match = True  
        for j in range(len(p)):           # Loop over characters  
            if t[i+j] != p[j]:            # compare characters  
                match = False             # mismatch; reject alignment  
                break  
        if match:  
            occurrences.append(i)          # all chars matched; record  
    return occurrences
```

P : word

T : There would have been a time for such a word


-----word-----word----->word
 -> -> ->

Exact matching: naïve algorithm

$$n = |P| \quad m = |T|$$

How many alignments are possible?

$$m - n + 1$$



Exact matching: naïve algorithm

$$n = |P| \quad m = |T|$$

Greatest # character comparisons possible?

$$n(m - n + 1)$$

P: aaaa

T: aa

aaaa aaaa aaaa aaaa aaaa aaaa aaaa aaaa aaaa

aaaa aaaa aaaa aaaa aaaa aaaa aaaa aaaa

aaaa aaaa aaaa aaaa aaaa aaaa aaaa aaaa

aaaa aaaa aaaa aaaa aaaa aaaa aaaa aaaa

aaaa aaaa aaaa aaaa aaaa aaaa aaaa aaaa

Exact matching: naïve algorithm

$$n = |P| \quad m = |T|$$

Least # character comparisons possible?

$$m - n + 1$$

P : abbb

T : bbb

abbb abbb abbb abbb abbb abbb abbb abbb abbb
abbb abbb abbb abbb abbb abbb abbb abbb
abbb abbb abbb abbb abbb abbb abbb abbb
abbb abbb abbb abbb abbb abbb abbb abbb
abbb abbb abbb abbb abbb abbb abbb abbb

Exact matching: naïve algorithm

How many character comparisons in this example?

P: word

T: There would have been a time for such a word

word word word word word word word word word word

word word word word word word word word word

word word word word word word word word word

word word word word word word word word word

word word word word word word word word word

Hint: there are 41 possible alignments

Exact matching: naïve algorithm

How many character comparisons in this example?

P: word

T: There would have been a time for such a word

word word word word word word word word word word

word word word word word word word word word word

word word word word word word word word word word

word word word word word word word word word word

word word word word word word word word word word

40 mismatches + 6 matches = 46 character comparisons

Closer to the minimum (41) than the maximum (164)

Exact matching: naïve algorithm

```
def naive(p, t):  
    occurrences = []  
    for i in range(len(t) - len(p) + 1):  # Loop over alignments  
        match = True  
        for j in range(len(p)):           # Loop over characters  
            if t[i+j] != p[j]:            # compare characters  
                match = False             # mismatch; reject alignment  
                break  
        if match:                          # all chars matched; record  
            occurrences.append(i)  
    return occurrences
```

Even more naïve:
remove **break**



Exact matching: better algorithms?

P: word

T: There would have been a time for such a word

.....word.....→
 →

u doesn't occur in *P*, so skip next two alignments

P: word

T: There would have been a time for such a word

.....word.....→
 word skip!
 word skip!
 word

We'll take such ideas further when we discuss Boyer-Moore